

Aviral Srivastava

• Github: [kebab-mai-haddi](#) • LinkedIn: [/in/sriavi](#) • Website: [aviralsrivastava.com](#)
• Email: hi@avisri.com • 617-283-3811

EDUCATION

Boston University, Boston, MA.

Sep 2019 - Dec 2020 (Expected)

- MS, Computer Science. CGPA : 3.77/4

Vellore Institute of Technology, Chennai, India.

Jul 2014 - May 2018

- B.Tech, Computer Science and Engineering. CGPA : 3.56/4

SKILLS

- **Languages:** Python, C++, PHP, Java, SQL
- **Technologies/Platforms:** Spark, Dask, Kubernetes, Docker, Vagrant, Argo, Airflow, Kafka, Pulsar, Celery, Flink, AWS, GCP, Azure, Snowflake, Cassandra, RocksDB, PrestoDB, KyotoDB, Redis, Django.

WORK

EXPERIENCE

Redhat, Data Engineering Intern, Internal Data-Hub team

May 2020 - Present

- Re-partitioning the schema of [SOS reports](#) to make data accessible.
- Building a data browsing platform to be used by AI and ML teams, using Spark, JupyterHub, Argo, Kubernetes, S3, and, Ceph.
- Built Argo workflows for AI and ML engineers to deploy their notebooks directly without any hidden technical debt such as configuration management, Spark OOM issues, etc.
- Reduced S3 object count of all data by more than 50 percent, removed the Ceph's limit on object count.

MIT, Research Associate (Software Engineering) @ ICORPS

Jun 2020 - Present

- I am working on commercializing my research work previously done at BU (see next work section).
- Leading the research team to conduct interviews for determining Customer Segments and Value Propositions.

Boston University, Research Associate (Software Engineering)

Sep 2019 - May 2020

- With [Dr Eric](#), I worked on making UML diagrams scalable and reliable for large and complex codebases.
- Stack: **Python3, Java, AST, OOP Design and Analysis**

Atlan, Data Engineer

May 2018 - Jul 2019

- Atlan is a data democratization startup that helps data teams collaborate frictionlessly on data projects.
- Ensured the development and on-premise deployments of dashboards like [United Nations SDGs monitoring](#), [Ujjwala Dashboard](#) using **Python3, Kubernetes, Docker and AWS S3**.
- Created serverless JupyterHub: a multi-user hub that spawns, manages, and proxies multiple instances of the single-user Jupyter notebook server. Stack: **papermill, Python3, Docker, Helm, Kubernetes**. Used by companies like **Unilever, Mindshare, McKinsey & Co.**
- Built and deployed Data Repository: a version control for data, including features like 'git diff', rollback, at a scale of over **billion rows in one cycle**. Stack: **Python3, S3, RocksDB, KyotoDB, Hive metastore and PrestoDB**. This scaled the production releases by **300%**: 1 project release to 3 projects release per quarter.
- Built a micro-service with Event Sourcing mechanism for pricing in [Collect](#) using **Kafka, Python3, Docker**.

Wingify, SDE Intern

Dec 2017 - Mar 2018

- Wingify is a B2B software product company with clients like **ebay, Disney, AMD, etc.**
- It's product is VWO which is an experience optimization platform that improves key business metrics by empowering clients to easily discover insights, test ideas, and improve engagement.
- As an intern, I worked with Head of Engg, Ankit Jain, to build a micro-service which parses textual inputs into Domain Specific Language trees and JavaScript code.
- APIs of this microservice scaled query processing of the product by **60%**.
- Received the full time employment offer based on my performance during internship.
- Stack: **PHP, MySQL, CICD, Docker, Vagrant, Blackbox Testing, Unit testing, JMeter**

Commutatus, SDE Intern

Jun 2016 - Jul 2016

- Worked on building CRM platforms on MVC architecture.

LeagueSX, SDE Intern

Dec 2015 - Feb 2016

- Python/Django Backend Developer Intern.

ENTR EXPERIENCE	Centify Co-founder, CTO	Apr 2017 - Apr 2018
	<ul style="list-style-type: none"> Centify was an audience Engagement tool which enabled digital media platforms to integrate social media reactions with video or textual content. We were selected in YC School with mentorship track, featured in IBM. 	
OPEN SOURCE	<ul style="list-style-type: none"> Profiling RocksDB-Flink Currently working with Dr Vasia Kalavri (Prime Member Contributor of Apache Flink) and Dr John Liagouris to determine overhead caused by state-ful operations in Flink streaming engine. ZCash Service Status Dashboard <ul style="list-style-type: none"> Zcash is a crypto-currency aimed at using cryptography to provide enhanced privacy for its users compared to other crypto-currencies like Bitcoin. Built a health-check service status dashboard to monitor the core ZCash blockchain, crypto exchanges, blockchain explorers and ZCash communities and forums. EEG Simulator <ul style="list-style-type: none"> Developed the first tool to simulate EEG readings using Kinetic LGN model with Dr Basab. Scaled the experimental analysis by more than 50%. Self Driving Remote Control Bot <ul style="list-style-type: none"> A collaborative autonomous driving bot to scale manual work in small-scaled factories. Footfall Tracker <ul style="list-style-type: none"> Identified eating patterns in campus dining to significantly reduce food waste by prediction. 	
HACKATHONS	<ul style="list-style-type: none"> Medhacks: Developed brain's state(fear, fright, flight) simulation. Won the verbal mention. Hackinout: Developed Chemical Component Detector in brain's thalamus. 	
RESEARCH EXPERIENCE	School of Computing Science , University of Manchester	Sep 2017 - Aug 2018
	<ul style="list-style-type: none"> Data Science Research Intern. 	
	Data Science Lab, Computer Science , University of Liverpool	May 2017 - Aug 2017
	<ul style="list-style-type: none"> Data Science Research Intern. 	