

Aviral Srivastava

• Github: [kebab-mai-haddi](#) • LinkedIn: [/in/sriavi](#) • Website: [aviralsrivastava.com](#)
• Email: hi@avisri.com • 617-283-3811

EDUCATION

Boston University, Boston, MA.

Sep 2019 - Dec 2020 (Expected)

- MS, Computer Science. CGPA : 3.77/4
 - **Subjects:** Designing Data Warehouse, Database Security, Advanced Database Management Systems, Graduate Distributed Systems (MIT), Machine Learning, Artificial Intelligence, Advanced Health Informatics and Big Data, Advanced Analysis of Algorithms, Information Systems Analysis and Design.

Vellore Institute of Technology, Chennai, India.

Jul 2014 - May 2018

- B.Tech, Computer Science and Engineering. CGPA : 3.56/4

SKILLS

- **Languages:** Python, C++, PHP, Java, Rust, SQL
- **Technologies:** Spark, Dask, Kubernetes, Docker, Vagrant, Argo, Airflow, Kafka, Pulsar, Celery, Flink, Snowflake, Cassandra, RocksDB, PrestoDB, KyotoDB, Redis, Django.
- **Platforms:** AWS, GCP, Azure

WORK

EXPERIENCE

Redhat, Data Engineering Intern, Internal Data-Hub team, Boston

May 2020 - Present

- Re-partitioned the schema of [SOS reports](#) (over 10 Trillion rows, 100+ TB data) to ensure high data availability and low latency. Reduced the count of S3 objects by more than 90%: from 15M+ objects to < 1M objects.
- Built a fault-tolerant, self-healing data browsing platform to be used by AI and ML teams w/o any hidden technical debts, using **Spark, JupyterHub, Argo, Kubernetes, S3**, and, **Ceph**.
- Built **Argo** workflows for AI and ML teams to deploy their adhoc Jupyter notebooks in automation mode.

Boston University, Research Associate (Software Engineering), Boston

Sep 2019 - May 2020

- With [Dr Eric](#), I worked on making UML diagrams scalable and reliable for large codebases.
- Got selected in **MIT-ICORPS** Fusion program to commercialize the research.
- Conducted customer discovery interviews with CEOs, CTOs, Engg Managers from various companies to understand the **Customer Segments** and **Value Propositions**.

Atlan, Data Engineer

May 2018 - Jul 2019

- Atlan is a data democratization startup that helps data teams collaborate frictionlessly on data projects.
- Ensured the development and on-premise deployments of dashboards like [United Nations SDGs monitoring](#), [Ujjwala Dashboard](#) using **Python3, Kubernetes, Docker** and **AWS S3**.
- Created serverless JupyterHub: a multi-user hub that spawns, manages, and proxies multiple instances of the single-user Jupyter notebook server. Stack: **papermill, Python3, Docker, Helm, Kubernetes**. Used by companies like **Unilever, Mindshare, McKinsey & Co.**
- Built and deployed Data Repository: a version control for data, including features like 'git diff', rollback, at a scale of over **billion rows in one cycle**. Stack: **Python3, S3, RocksDB, KyotoDB, Hive metastore** and **PrestoDB**. This scaled the production releases by **300%**: 1 project release to 3 projects release per quarter.
- Built a micro-service with Event Sourcing mechanism for payments in [Collect](#) using **Kafka, Python3, Docker**.

Wingify, SDE Intern

Dec 2017 - Mar 2018

- Wingify is a B2B software product company with clients like **ebay, Disney, AMD, etc.**
- It's product is VWO which is an experience optimization platform that improves key business metrics by empowering clients to easily discover insights, test ideas, and improve engagement.
- As an intern, I worked with Head of Engg, Ankit Jain, to build a micro-service which parses textual inputs into Domain Specific Language trees and JavaScript code.
- APIs of this microservice scaled query processing of the product by **60%**.
- Stack: **PHP, MySQL, CICD, Docker, Vagrant, Blackbox Testing, Unit testing, JMeter**

Other Internships

- Commutatus, LeagueSX, University of Manchester, University of Liverpool

OPEN SOURCE

- [Profiling RocksDB-Apache Flink](#)
- [ZCash Service Status Dashboard](#)
- [Other projects](#)